

Consumer Decision-Making Analysis in R

Reproducible, Collaborative, and AI-Aware Workflows

Who we are



Jannis
Kreienkamp

*Statistics
Software*



Maximilian
Agostini

*Data OPS
Infrastructure*

Why this workshop?



Overview

Structure of the workshop

- Day 1: The scaffold
- Day 2: R and data analysis
- Day 3: Producing a report + Supercharge

The three-day roadmap

Day 1: Scaffold

- R
- Rstudio
- Git/GitHub basics
- Folders
- Data protection
- Quarto

Day 2: Analyse

- Base R
- Import
- Clean
- Transform
- Visualise
- Describe
- Model
- Moderate / mediate
- Collaborate

Day 3: Publish

- Export tables and figures
- Write methods
- Write results
- Static/ Annotated/ interactive report
- AI-use statement
- Reproducibility check

Day 1

The Repository

What, where, why?

A quick vocabulary check

fork	commit	push / pull	branch / merge	stage
repository	RStudio	.Rproj	renv	snapshot
Quarto	render	.qmd	GitHub Pages	Overleaf


Don't know them yet? Perfect. After today, we hope you do.


Setting expectations — not today

Deep R syntax	→ Day 2
VS Code	→ Day 3
Real data	→ Day 2/3
Responsible AI use	→ Day 3

Today's single job: install the tools and
understand how they fit together.

Good Scientist Recipe







kubernetes

conda



The L^AT_EX Project



Overleaf



stackoverflow



Posit Community

StackExchange



HTML



AS PREDICTED



OSF REGISTRIES



python™



julia



R Studio®



Visual Studio Code

Windsurf

CRAN



OSF

The Dataverse® Project

zenodo



git



GitHub



GitLab



Hugging Face




LM Studio




Claude



Continue





Every tool today is free and open source (FOSS)

R & friends	free; huge community; you keep your analysis (vs SPSS: €€€, closed)
Markdown	plain text; machine-readable; no vendor can lock or change it (vs Word)
Git / GitHub	open standard; the full history is yours to keep and to move
Zotero	free, open reference manager; plugs into Word, Quarto, Overleaf
Quarto	free, open publishing; one source → website, PDF, Word

Open + free + standard = reproducible, shareable, and future-proof.


No licence to expire. No format you can't open in ten years.


the-science-repository


Public

A template for reproducible science in R: one repo for your data, analysis, interactive render, and manuscript.

 R  5







A quick tour of the folders (the kitchen)

the-science-repository/

data/

INGREDIENTS



R/

APPLIANCES & TECHNIQUES



reports/

THE RECIPES



renders/

THE PLATED DISHES



references/

THE NUTRITION FACTS



A quick tour of the folders (the kitchen)

the-science-repository/	
└─ README.md	Instructions in every folder
└─ data/	INGREDIENTS
└─ README.md	
└─ mock/	← display/sample ingredients (safe to share, committed)
└─ raw/	← ingredients as delivered (private, gitignored)
└─ processed/	← prepped ingredients (private, reproducible)
└─ R/	APPLIANCES & TECHNIQUES
└─ README.md	
└─ functions/	← reusable base techniques (load, model, plot)
└─ reports/	THE RECIPES
└─ webpage/	← recipe for the website dish
└─ manuscript/	← recipe for the paper dish
└─ renders/	THE PLATED DISHES
└─ webpage/	← the served website (100% generated)
└─ manuscript/	← the served paper (mixed: hand- + machine-plated)
└─ generated/	← machine-plated parts (figures, tables, methods/results)
└─ sections/	← hand-plated parts (abstract, intro, discussion)
└─ references/	THE NUTRITION FACTS (inspiration/ shared references.bib)

A quick tour of the folders (the kitchen)

the-science-repository/

— README.md

Instructions in every folder

— data/

INGREDIENTS

— README.md

— mock/

← display/sample ingredients (safe to share, committed)

— raw/

← ingredients as delivered (private, gitignored)

— processed/

← prepped ingredients (private, reproducible)

— R/

APPLIANCES & TECHNIQUES

— README.md

— functions/

← reusable base techniques (load, model, plot)

— reports/

THE RECIPES

— webpage/

← recipe for the website dish

— manuscript/

← recipe for the paper dish

— renders/

THE PLATED DISHES

— webpage/

← the served website (100% generated)

— manuscript/

← the served paper (mixed: hand- + machine-plated)

— generated/

← machine-plated parts (figures, tables, methods/results)

— sections/

← hand-plated parts (abstract, intro, discussion)

— references/

THE NUTRITION FACTS (inspiration/ shared references.bib)

A quick tour of the folders (the kitchen)

the-science-repository/

— README.md

Instructions in every folder

— data/

— README.md

— mock/

— README.md

— raw/

— README.md

— processed/

— README.md

— R/

— README.md

— functions/

— README.md

— reports/

— webpage/

— README.md

— manuscript/

— README.md

— ...

Every project explains itself

A README tells you:

- where you are in the project
- what lives in the folder
- how to set up and run things



READMEs are written in Markdown

```
# Heading level 1
## Heading level 2

**bold**, *italic*, ~~strikethrough~~

- bullet point
- another bullet

[My link label](https://example.com)

---

```R
multiple lines of
code instructions
```

My `variable name` looks like code

:+1:
```

Heading level 1

Heading level 2

bold, *italic*, ~~strikethrough~~

- bullet point
- another bullet

[My link label](#)

```
# multiple lines of
# code instructions
```

My `variable name` looks like code



READMEs are written in Markdown

Why Markdown instead of Word?

A README written in Markdown is still readable even before it is rendered.

Git-friendly

- Plain text → every change is visible
- Works naturally with version control (``git diff``)

Predictable

> What you write is what everyone sees.

- No hidden formatting
- No accidental style changes
- No "Why did Word move everything again?"

Future-proof

- Open format → no vendor lock-in
- Readable in any text editor
- Likely to remain accessible for decades
- Readable by programs and AI

Why Markdown instead of Word?

A README written in Markdown is still readable even before it is rendered.

Git-friendly

- Plain text → every change is visible
- Works naturally with version control (`git diff`)


Predictable

What you write is what everyone sees.

- No hidden formatting
- No accidental style changes
- No "Why did Word move everything again?"

Future-proof


- Open format → no vendor lock-in
- Readable in any text editor
- Likely to remain accessible for decades
- Readable by programs and AI



Version control: the safety net



git



GitHub



Why should I care

Your personal work milestone timeline




Why should I care

Your personal work milestone timeline



Your work



Collaborator work




Why should I care

Your personal work milestone timeline



Your work




Collaborator work



Why should I care

Your personal work milestone timeline



The words you'll hear all day

| | |
|--------|---|
| Fork | Make your own copy of someone else's repo on GitHub |
| Clone | Download a repo to your computer |
| Stage | Pick which changes go into the next commit |
| Commit | Save a labelled snapshot of the staged changes |
| Push | Send your commits up to GitHub |
| Pull | Bring GitHub's changes down to your computer |
| Branch | A separate line of work |
| Merge | Combine work from different branches |

Errors are part of the process



Errors are normal. Everyone in this room will see red text today.
An error message is a clue, not a failure.

When you hit one:

1. Read it slowly – the answer is usually right there (a path, a name, a missing thing)
2. Recall the last thing you did – the error points back to it
3. Copy the exact message into a search engine – someone has hit it before
4. Ask a neighbour, or ask us

Step 1 — get an account

1. Go to `github.com`
2. Sign up (use an email you'll keep – ideally your university email)
3. Verify your email
4. Pick a username you're happy to share publicly

CHECK

Can you log in and see your (empty) GitHub dashboard?

WATCH OUT

- Verification email can be slow or land in spam – check both.
- Some corporate/university mail servers block the verification link (check the spam)
- The username becomes part of your public website URL later – choose something you won't regret printing on a paper.

Step 2 — make it yours

1. Open: `github.com/thedataflowcompany/the-science-repository`
2. Click "Fork" (top right)
3. Keep the name, click "Create fork"
4. You now have: `github.com/<your-awesome-username>/the-science-repository`

WATCH OUT

- If you've forked this repo before, GitHub won't fork it twice – just open your existing copy.
- Don't accidentally fork into an organisation; fork into your own account.

LATER — STARTING YOUR OWN PROJECT (NOT TODAY)

Forking keeps a visible link to this template – ideal for the workshop.

For your own real project, start a clean, independent repo instead:

1. Code → Download ZIP (a snapshot with no fork link)
2. Create a NEW repo on GitHub – public, or private if your data or project demands it
3. Copy the files in, commit, and push

Today we all stay on the public fork.

Step 3 — publish your analysis website

```
Preview Code Blame 87 lines (66 loc) · 5.02 KB
1  ✓ # The Science Repository
2
3  A template for a modern, reproducible science project in R. One git repository
4  holds the data, the analysis engine, and two polished outputs built from
5  it: a website that shows the whole analysis transparently, and a journal
6  manuscript. Fork it, rename it, and fill it with your own work.
7
8  > 📄 [Open the full analysis website →](https://thedataflowcompany.github.io/the-science-repository/renderers/webpage/)
9  > Every cleaning step, figure, table, and model, rendered from the code in this repo.
10
11  ---
12
13  ✓ ## What lives where
14
15  Start here, then click into any folder's own README for the details.
16
17  | Folder | What it holds | More |
18  | --- | --- | --- |
19  | [R/](R/) | The analysis engine: setup + reusable functions. The single source of truth for the numbers. Nothing runs on its own
```

Step 3 — publish your analysis website

The Science Repository

A reproducible consumer decision-making analysis in R — workshop edition.

[View the Project on GitHub](#)

TheDataFlowCompany/the-science-repository

The Science Repository

A template for a modern, reproducible science project in R. One git repository holds the **data**, the **analysis engine**, and two polished outputs built from it: a **website** that shows the whole analysis transparently, and a **journal manuscript**. Fork it, rename it, and fill it with your own work.

 [Open the full analysis website](#) → Every cleaning step, figure, table, and model, rendered from the code in this repo.

What lives where

Start here, then click into any folder's own README for the details.

| Folder | What it holds | More |
|--------|---|-----------------------------|
| R/ | The analysis engine : setup + reusable functions. The single source of truth for the numbers. Nothing runs on its own — the reports call it. | R/README.md |

Step 3 — publish your analysis website

In YOUR fork on GitHub:

1. Settings → Pages
2. Source: "Deploy from a branch"
3. Branch: main • Folder: / (root) → Save
4. Wait ~1 minute, then open:
<https://<your-username>.github.io/the-science-repository>
5. While waiting, apply for GitHub Education

WATCH OUT

- GitHub Pages on a FREE account needs a PUBLIC repo. Your fork of a public template is public by default, so you're fine. A private repo needs GitHub Pro / Education for Pages.
- First build takes 1–2 minutes. A 404 usually means the build isn't done yet, or the URL is wrong – it must end with `/renders/webpage/` .

Step 4 — edit the README link and commit


The README points to the project's website. Make it point to YOURS.

1. Open README.md on GitHub → click the pencil (edit)
2. Find the website link near the top
3. Replace the URL with your own github.io address (from the previous step)
4. Scroll down → "Commit changes"

Message: "Point website link to my fork" or anything else you like

What just happened, visually

TheDataFlowCompany/the-science-repository



The words you'll hear all day

- ✓ Fork Make your own copy of someone else's repo on GitHub
- Clone Download a repo to your computer
- Stage Pick which changes go into the next commit
- ✓ Commit Save a labelled snapshot of the staged changes
- Push Send your commits up to GitHub
- Pull Bring GitHub's changes down to your computer
- Branch A separate line of work
- Merge Combine work from different branches

Data protection checkpoint .gitignore

Usually safe to commit:

- scripts / code
- README & docs
- mock (synthetic) data
- rendered outputs

Usually NOT safe to commit:

- real / identifiable participant data
- passwords, tokens, API keys
- ethics docs with sensitive details
- anything confidential

Rule of thumb:

Scripts can be shared. Raw data usually cannot.
Outputs depend on whether anyone can be identified.

Real data can't be committed by accident

| | | |
|-----------------|------------|--|
| data/mock/ | committed | (synthetic, safe – everyone uses this today) |
| data/raw/ | gitignored | (your real data – never leaves your machine) |
| data/processed/ | gitignored | (cleaned data – also stays local) |

.gitignore also blocks common data file types everywhere except mock/.

BREAK


Coffee break — 10:30–11:00

Back at 11:00.

✓ So far: GitHub account • your fork • website live on Pages •
first commit

→ Next: install R & RStudio, then clone your fork to your own
laptop

The engine and the cockpit



The engine and the cockpit

R = the programming language that does the work → the stove
RStudio = the environment you work in: write code, see data,
run scripts, view plots, manage the project (the cockpit)

You can run R without RStudio – but RStudio makes everything visible.

Step 5 — install the software

Go to: <https://posit.co/download/rstudio-desktop> (has links to both)

1. Install R: cran.r-project.org (pick your OS, latest version)
2. Install RStudio: posit.co/download/rstudio-desktop (free Desktop)
3. Open RStudio – it finds R automatically

CHECK

Does RStudio open without errors?

In the Console, type `1 + 1` and press Enter → you should see `[1] 2`

WATCH OUT

- Install R FIRST, then RStudio (RStudio needs R to exist).
- University laptops often need admin rights
- Mac: pick the build for your chip (Apple Silicon vs Intel).
- Antivirus on Windows can quarantine the installer – allow it.

Step 6 — clone your fork (1 / 2)

Cloning = downloading YOUR fork (with its full history) to your machine.

In RStudio:

File → New Project → Version Control → Git

Repository URL: <https://github.com/<your-username>/the-science-repository> → Create Project

Two ways to handle the GitHub login – pick one:

A. GitHub Desktop (easiest today)

Install desktop.github.com → sign in (in the browser) → File → Clone

No Personal Access Token to set up.

B. Personal Access Token / PAT (the sustainable, long-term setup)

In the R Console:

```
usethis::create_github_token() # opens GitHub to create the token
```

```
gitcreds::gitcreds_set()      # paste it once; RStudio remembers it
```

Alternative:

Go to Github/Settings/Developer Settings/Personal Access Tokens/PAT

Heads up: this is the fiddliest step of the whole day. That's expected – we will get everyone through it. Take a breath; ask for help.

Step 6 — clone your fork (2 / 2)

WATCH OUT

- If the "Git" option is greyed out, Git isn't installed → install from git-scm.com, restart RStudio.
- FIRST TIME ONLY – tell Git who you are (same email as GitHub):

```
git config --global user.name "Your Name"
git config --global user.email "you@university.edu"
```
- Pushing over HTTPS no longer accepts your password. When asked, paste a Personal Access Token (option B), or use GitHub Desktop (option A) which signs you in through the browser and handles credentials for you.

Always open via the .Rproj

the-science-repository.Rproj  **x2**

Opening the project (not single files) tells R:

- where the project root is
- where data, scripts, and outputs live
- file paths work the same on everyone's computer

How RStudio Projects get created

File → New Project → ...

1. New Directory a brand-new empty project folder
2. Existing Directory wrap a project around a folder you already have
3. Version Control clone from GitHub (what we just did)

Every option produces a `.Rproj` file = the project's anchor.

BREAK


Lunch — 12:30–13:30

Back at 13:30.

✓ So far: R & RStudio installed · fork cloned · project open
via the .Rproj

→ Next: restore the exact packages, then the publishing
pipeline and your
first local render

Renv: pinning package versions



Why pinning package versions matters



Same code + different package versions = different (or broken) results.

"It works on my laptop but not yours."
"It worked last year but not today."

renv = a locked, per-project package library

| | |
|-----------|---|
| renv.lock | records the exact version of every package |
| renv/ | a private package library just for this project |
| .Rprofile | activates renv when the project opens |

→ Everyone restores the SAME versions. No more version drift.

Step 7 — install the exact packages

```
install.packages("renv") # only if RStudio doesn't prompt you
renv::restore()          # installs every package at the pinned version
```

R


CHECK

renv reads renv.lock and installs what's missing.
This can take a few minutes – that's normal.
When it finishes, you have every package the project needs.

WATCH OUT

- If prompted "Do you want to proceed? [y/N]" – type y and Enter.
- First restore can compile packages from source. Mac users may be asked to install the Xcode Command Line Tools – accept it.
- It can take 5–15 min on a fresh machine. Let it finish; don't interrupt.
- Needs a stable internet connection.

Quarto: Rendering outputs



A report is a recipe, read top to bottom


reports/webpage/

| | | |
|-------------------------|-----------------------|------------------------------|
| 01-data-preparation.qmd | load + clean the data | → writes the processed cache |
| 02-descriptives.qmd | summarise it | → uses the processed data |
| 03-analysis.qmd | model it | → uses the processed data |

Each page runs top → bottom. Setup first, results last.

They all call the shared functions in R/functions/ (the kitchen techniques).


Data Folders




raw → processed → mock (what OL-data-preparation handles)

| | |
|-----------------|---|
| data/raw/ | original data, exactly as received – never hand-edited (private) |
| data/processed/ | cleaned, analysis-ready – created BY a script, can be rebuilt (private) |
| data/mock/ | synthetic stand-in – committed, safe to share, used in the workshop |

Zotero



Manage references once, cite everywhere




Why Zotero? Free and open (vs paid managers), and it plugs into Word, Quarto, and Overleaf. One library, cited everywhere.

Website and manuscript, same numbers

| | | | |
|---------------------|-----------------|---------------------|----------------------|
| reports/webpage/ | —quarto render→ | renders/webpage/ | (HTML website) |
| reports/manuscript/ | —quarto render→ | renders/manuscript/ | (LaTeX / Word paper) |

Both reuse R/. Same data, same functions, two presentations.

Rendering a webpage



Step 8 — build the site locally

```
quarto render reports/webpage
```

BASH

→ rebuilds renders/webpage/
Open renders/webpage/index.html to see it.
(Quarto ships inside RStudio – no separate install needed.)

WATCH OUT

- Run this in the TERMINAL tab, not the R Console.
- Render the WHOLE site (reports/webpage), not a single page – the first page builds the data/processed cache the others depend on.
- If "quarto: command not found", update RStudio (recent versions bundle Quarto) or install from quarto.org.


BREAK

Coffee break — 15:00–15:30

Back at 15:30.

✓ So far: environment restored • website rendered locally –
the Day 1 finish line
→ Next: the optional second output (the PDF manuscript), common
blockers, and
the end-of-day checklist

Rendering a manuscript (if there is time)



The second output — the paper, and how co-authors join in

The website is today's finish line. The manuscript is the SAME engine, second dish.

To build the PDF you need a LaTeX engine (one-time install):

in the R Console:

```
install.packages("tinytex")  
tinytex::install_tinytex()
```

Then:

```
quarto render reports/manuscript  
→ generates the data-driven parts (methods, results, figures, tables)  
→ renders/manuscript/main.tex assembles them into manuscript.pdf
```

Overleaf:

the self-contained renders/manuscript/ folder can sync to Overleaf
→ online, collaborative LaTeX writing with co-authors

WATCH OUT

- tinytex downloads a full LaTeX distribution – needs internet and a few minutes.
- The manuscript needs the LaTeX engine above; the WEBSITE never does.

The day's usual suspects (and the fix)

| | |
|---------------------------------|--|
| "Git option greyed out" | → install Git (git-scm.com), restart RStudio |
| "Authentication failed" on push | → use a Personal Access Token, not a password |
| "Could not find function" | → package not loaded → did <code>renv::restore()</code> finish? |
| "cannot open file ... raw/..." | → you're in real mode; default is mock (do nothing) |
| Website shows 404 | → wait for the build; URL ends <code>/renders/webpage/</code> |
| "quarto: command not found" | → update RStudio / install from quarto.org |
| A single page won't render | → render the whole <code>reports/webpage</code> folder |

Before you leave today

- [✓] GitHub account created
- [✓] Repository forked to your account
- [✓] GitHub Pages on → your website loads
- [✓] First commit made (README link points to your site)
- [✓] R and RStudio installed
- [✓] Fork cloned and opened via the .Rproj
- [✓] `renv::restore()` finished successfully
- [✓] quarto render reports/webpage works locally



Achievement unlocked
Software Developer Lvl 1

Day 2 — analyse the data in R

tomorrow we:

- work inside R/ and the report pages
- import, clean, and inspect the data
- build descriptives, figures, and models
- use branches and merges to collaborate in pairs

See you tomorrow

We're happily staying for extra questions

info@thedataflowcompany.com
www.thedataflowcompany.com

